

LLM Load Test Report

Model: Qwen/Qwen3.8-27B-FP8

Server: http://localhost:8000

Total rounds: 6

Input token range: 2000 - 2000 (step 2000)

User levels: [10, 20, 40, 60, 80, 100]

Max output tokens: 2000

Temperature: 0.7


Generated: 2026-09-17 13:27:05

GPU before: 1, [N/A], [N/A], 11.92


GPU after: 96, [N/A], [N/A], 24.37

Aggregate Output Throughput vs Input Tokens


Aggregate Output Throughput vs Input Tokens




Input Processing Throughput vs Input Tokens




p50 Latency vs Input Tokens




Request Success Rate vs Input Tokens




Output Throughput Per User vs Input Tokens



Average Per-Request Output Speed vs Input Tokens




p90 Latency vs Input Tokens




p99 Latency vs Input Tokens


p99 Latency vs Input Tokens



Aggregate Output Throughput vs Concurrent Users



p50 Latency vs Concurrent Users



Detailed Results

Round	Users	Input	OK	Fail	Success %	Actual Input	Output	Input TPS	Output TPS	TPS/User	p50	p95	p99	Max
1	10	2000	10	0	100.0%	1363	20000	45.8	67.1	6.7	297.74	297.87	297.87	297.87
2	20	2000	20	0	100.0%	1363	40000	79.3	116.4	5.8	343.24	343.54	343.68	343.68
3	40	2000	40	0	100.0%	1363	80000	124.5	182.7	4.6	436.94	437.86	437.86	437.86
4	60	2000	60	0	100.0%	1363	120000	154.4	226.6	3.8	527.81	529.35	529.48	529.48
5	80	2000	80	0	100.0%	1363	160000	171.8	252.1	3.2	632.29	634.29	634.43	634.56
6	100	2000	100	0	100.0%	1363	200000	187.0	274.4	2.7	725.63	728.60	728.73	728.73