

LLM Load Test Report

Model: Qwen/Qwen3.8-27B-FP8

Server: http://localhost:8000

Total rounds: 4

Input token range: 7000 - 7000 (step 7000)

User levels: [100, 120, 140, 150]

Max output tokens: 2000

Temperature: 0.7


Generated: 2026-09-17 16:01:43

GPU before: 0, [N/A], [N/A], 11.90


GPU after: 96, [N/A], [N/A], 32.44

Aggregate Output Throughput vs Input Tokens

Aggregate Output Throughput vs Input Tokens




Input Processing Throughput vs Input Tokens




p50 Latency vs Input Tokens


p50 Latency vs Input Tokens




Request Success Rate vs Input Tokens



Output Throughput Per User vs Input Tokens




Average Per-Request Output Speed vs Input Tokens




p90 Latency vs Input Tokens

p90 Latency vs Input Tokens




p99 Latency vs Input Tokens


p99 Latency vs Input Tokens



Aggregate Output Throughput vs Concurrent Users



p50 Latency vs Concurrent Users



Detailed Results

Round	Users	Input	OK	Fail	Success %	Actual Input	Output	Input TPS	Output TPS	TPS/User	p50	p95	p99	Max
1	100	7000	100	0	100.0%	4594	200000	600.3	261.3	2.6	760.42	764.84	765.10	765.23
2	120	7000	120	0	100.0%	4594	240000	630.4	274.4	2.3	868.34	874.08	874.49	874.49
3	140	7000	0	140	0.0%	N/A	0	0.0	0.0	0.0	N/A	N/A	N/A	N/A
4	150	7000	0	150	0.0%	N/A	0	0.0	0.0	0.0	N/A	N/A	N/A	N/A

Sample Errors

Round 3: 7000 input tokens / 140 users

timed out

Round 4: 7000 input tokens / 150 users

timed out